

Basics of Computing

Peter Ruprecht

peter.ruprecht@colorado.edu

www.rc.colorado.edu

Slides:

https://github.com/ResearchComputing/Basics_Supercomputing

Outline

- Hardware hierarchy: supercomputer to core
- Components of a compute node
- Interconnect
- Storage and filesystems
- Coprocessors
- Operating systems

Compute Hardware Architecture

Processing: Supercomputer – Node – Socket/CPU – Core

Memory: Distributed – RAM – L3 – L2 – L1

Interconnect network

Storage (disk)

Coprocessors

Non-cluster supercomputers (IBM Blue Gene, SGI UV)

Cloud?

Parallelism at all levels of the computing system is the name of the game today ... SIMD, OpenMP, MPI

Clusters

Stampede (TACC)



6400 nodes

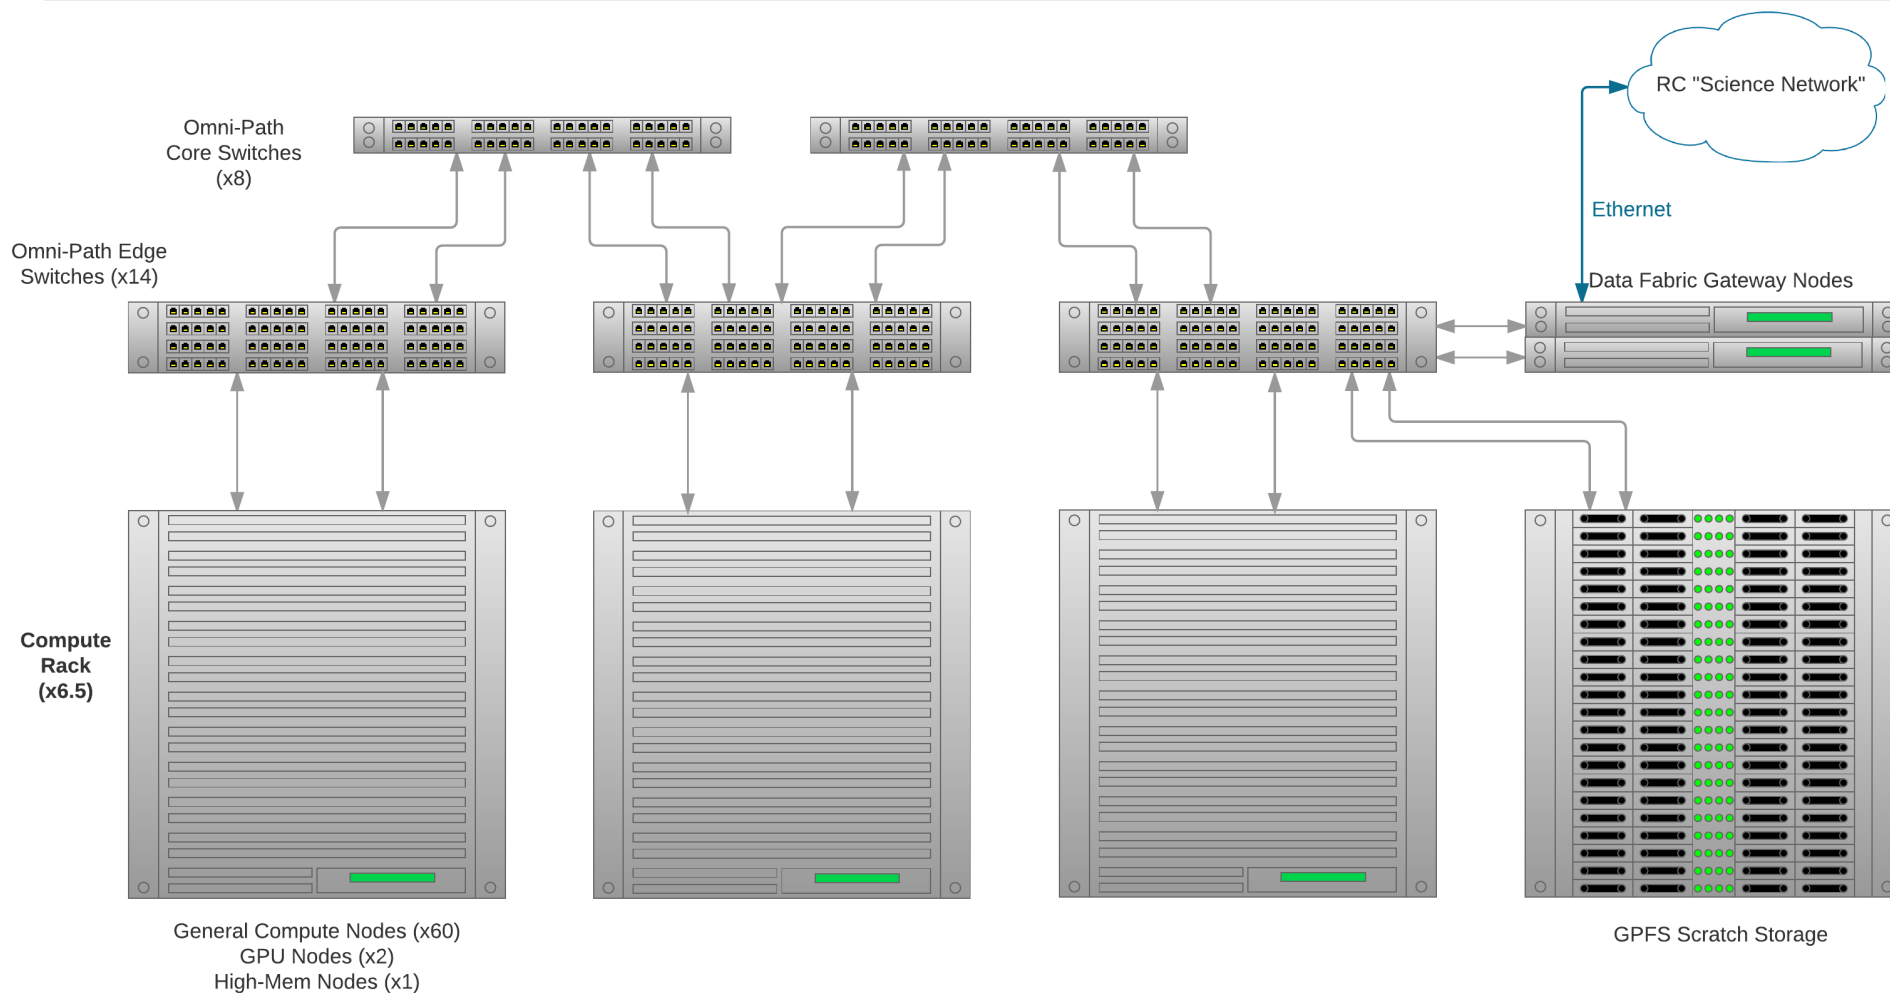
2 sockets per node, 8 cores per socket

56 Gbps Infiniband interconnect

14 PB parallel storage system


SUMMIT SCHEMATIC

Peter Ruprecht | July 15, 2016




Nodes and Processors

CPU sockets



Memory slots



Images: Supermicro, Intel

Host Processor Types

- x86_64 (Intel, AMD)
 - Most likely to encounter
 - General purpose
- Xeon Phi (Intel)
- POWER (IBM)
 - May have tight integration with coprocessor
- ARM
 - RISC, means fewer transistors
 - Low power

All processors operate at a certain frequency (several GHz) and can perform several instructions per cycle.

Computer Bus Layout

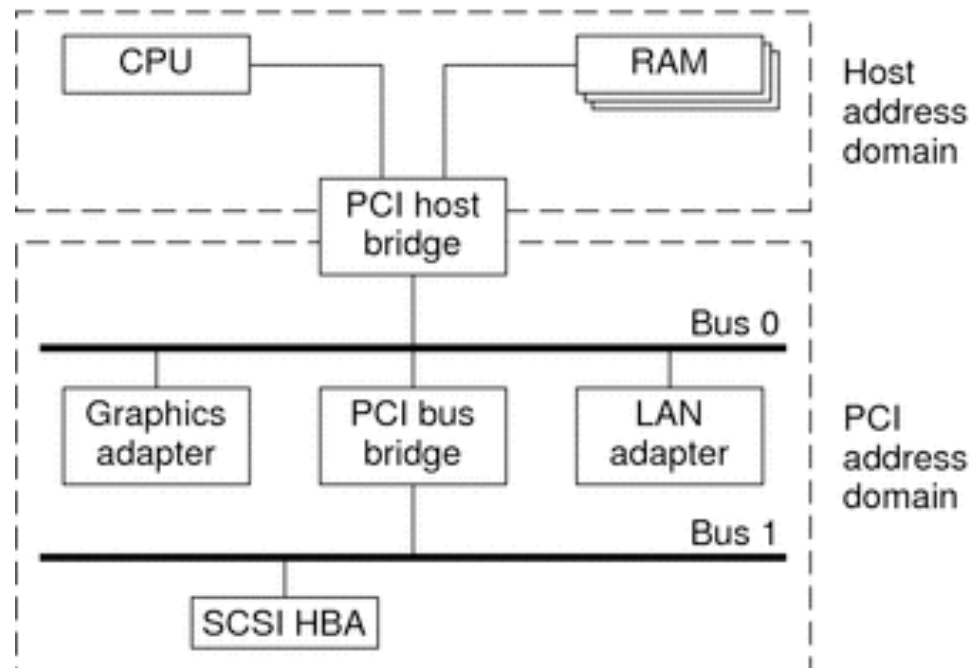


Image: Oracle



Memory

- Holds data that is being calculated on, as well as computational instructions
- Closest memory to CPU is actually on the processor die
 - Cache: Level 1 and 2 are dedicated to a single core
 - Level 3 is larger and shared between all cores in a socket
- Next closest is RAM
- Can also access memory on other nodes, via RDMA (remote direct memory access)

Shared memory is local to one node and several process threads can share the same data addresses.

Distributed memory is on multiple nodes and each process normally has its own copy or part of the data.

Shared and Distributed Memory



Interconnect

- How to access distributed memory across nodes?
- Need a fast, *low-latency* network.
- Current interconnect technologies
 - InfiniBand (40, 56, 100 Gbps)
 - Ethernet (10, 40, 100 Gbps)
 - Aries/Gemini (Cray)
 - OmniPath (100 Gbps)
- RDMA normally requires special interconnect-aware libraries; frequently these are included with MPI

Interconnect is what makes a bunch of nodes into a supercomputer!

Storage

Different types of “disk” for different needs

- Local disk in the node, often SSD
- Shared scratch
 - Often accessed over cluster interconnect
 - Parallel filesystems, eg Lustre or GPFS
 - Traditionally tuned for high bandwidth, not high IOPS
 - May have a “burst buffer” layer in front of it
 - Short-term storage only!!
- Longer-term or archive
 - Often uses Hierarchical Storage Management

CPU pipeline



Optimizing for Data Access

- Page Fault, file on IDE disk: 1,000,000,000 cycles
- Page Fault, file in buffer cache: 10,000 cycles
- Page Fault, file on ram disk: 5,000 cycles
- Page Fault, zero page: 3,000 cycles
- Main memory access: about 200 cycles
- L3 cache hit: about 52 cycles
- L1 cache hit: 2 cycles

The Core i7 can issue 4 instructions per cycle. So a penalty of 2 cycles for L1 memory access means a missed opportunity for 7 instructions.

Co-Processors

Specialized, usually massively multi-core, separate from host processor (CPU).

- GPU (Graphics Processing Unit, “video card”)
 - Thousands of cores
 - Great for vectorizable or embarrassingly-parallel apps
 - Programming frameworks include CUDA, OpenACC, OpenCL
 - Many applications (eg, Matlab) support CUDA directly
- Xeon Phi (being phased out)
 - Dozens of cores
 - Existing x86 code can be directly recompiled to run on Phi
 - Function offload

Operating Systems in HPC

Operating system: software that manages computer hardware and the processes running on it.

- MS Windows – hasn't really gotten a foothold
- Mac OSX – even less of a factor
- Unix / Linux – some variation will be on virtually any HPC system you use
 - May be a stripped-down version optimized for the particular hardware, or
 - May be a full distribution

What is Linux?

- Part of the Unix family of operating systems.
- Started in early '90s by Linus Torvalds.
- Technically refers only to the kernel; software from the GNU project and elsewhere is layered on top to form a complete OS. Most is open source.
- Several full distributions are available – from enterprise-grade, like RHEL or SUSE, to more consumer-focused, like Ubuntu.
- Runs on everything from embedded systems to supercomputers.

Why use Linux?

- Linux command-line syntax may seem overwhelming to the new user, but:
- It's the default operating system on virtually all HPC systems
- It's extremely flexible
- It tries not to get in your way
- It's fast and powerful
- Open-source scientific applications are developed predominantly for Linux
- You can get started with a few basic commands and build from there

Alternatives to the Command Line

- Domain-specific portals
 - CIPRES
 - SEAgrid
 - Many, many others
- Notebooks
 - Jupyter (“Julia, Python, R” , but many others)
- Domain-independent portal environments
 - Sandstone HPC
 - Viewpoint

users

shell: bash, csh

programs utilities

Linux kernel

Computer hardware

References

1. Eijkhout, Victor. *Introduction to High-Performance Scientific Computing*, 2015.
<https://bitbucket.org/VictorEijkhout/hpc-book-and-course/src>
2. Hager, G, and G Wellein. *Introduction to High Performance Computing for Scientists and Engineers*, CRC Press, 2010.
3. Levesque, John, and Gene Wagenbreth. *High Performance Computing*. CRC Press, 2010.
4. Neeman, Henry. *Supercomputing in Plain English*, High Performance Computing Workshop Series, <http://www.oscer.ou.edu/education.php>

Questions?

- Email rc-help@colorado.edu
- Link to survey on this topic:
<http://tinyurl.com/curc-survey16>

Speaker: Pete Ruprecht

Title: Computing Basics (July 2017 BSW)

- Slides:
https://github.com/ResearchComputing/Basics_Supercomputing